

Bala Sai Srikar Mandava

[✉ E - mail](#) | [📱 Mobile](#) | [🌐 LinkedIn](#) | [🐙 GitHub](#) | [📁 Portfolio](#)

SUMMARY

Data Engineer experienced in building and optimizing large-scale ETL pipelines, cloud data infrastructure, and analytics workflows. Proficient in Python, PySpark, Azure Data Factory, Databricks, dbt, and Snowflake with knowledge in incremental loading and pipeline reliability, processing 1.4Cr+ records across high-volume enterprise systems.

EDUCATION

SRM University AP

2021-2025

Computer Science and Engineering

8.84

- *Relevant Coursework:* Data Structures, Algorithms, OOPs, Cloud

WORK EXPERIENCE

LTM

Aug 2025 - Now

Data Engineer

Bangalore

- Redesigned ETL pipelines processing 1.4Cr+ records into parallel execution architectures using workload partitioning, async API processing, and incremental load strategies for improved scalability and reduced processing time.
- Enhanced pipeline reliability for high-volume data ingestion by implementing adaptive throttling, retry mechanisms, automated validation, and robust exception handling across Snowflake-integrated data pipelines.
- Streamlined batch-processing workflows with automated execution tracking, file-ingestion monitoring, dependency handling, and email-based reporting for distributed data systems.

PROJECTS

Retail Sales Analytics with Dimension Tracking

Azure Data Factory — Azure Databricks — PySpark — Azure Storage — Azure Key Vault — Power BI

- Developed an end-to-end retail analytics pipeline ingesting and transforming large-scale sales datasets using Azure Data Factory for orchestration and Azure Databricks for distributed processing.
- Engineered PySpark-based ETL workflows implementing Star Schema dimensional modeling, data cleansing, and validation with enterprise-grade security via Azure Key Vault and Storage Accounts.
- Developed Power BI dashboards integrated with Azure Logic Apps orchestration to deliver real-time sales performance insights and KPI tracking for business stakeholders.

Fuel Price Analytics Pipeline

[GitHub](#) | [Live](#)

Python — PostgreSQL — dbt — GitHub Actions — Streamlit

- Designed a zero-infrastructure serverless ETL pipeline using GitHub Actions CI/CD workflows to automate daily fuel price data ingestion and processing across 97+ Indian cities with zero hosting cost.
- Developed async Python web scraping workflows with transaction-safe bulk upsert operations in PostgreSQL and dbt Core transformation models for pricing aggregation, trend analysis, and 26+ data quality tests.
- Delivered real-time business intelligence via an interactive Streamlit-Plotly dashboard with motorized HTML email reporting, optimized through concurrent task execution and fault-tolerant retry mechanisms.

SKILLS

Languages

Python, Bash, MySQL

Data Engineering

Apache Airflow, Azure Data Factory, Azure Databricks, PySpark, dbt, Snowflake Cortex

Tools

Git, Power BI, Streamlit

Cloud

Azure (ADF, Databricks, Storage, Key Vault), AWS, Snowflake

CERTIFICATIONS

- Postman API Fundamentals Student Expert
- OCI Foundations Associate Certified
- AWS Educate Member